Languages, Compilers, Options and Astronomical Images

A 2026 reprise of Shortridge (2019), with sixteen new languages

Karl Glazebrook*
(with assistance from Claude Sonnet 4.7 by Anthropic)

April 2026

Abstract

This report extends Keith Shortridge's 2019 ADASS study (paper P9-17, `KnaveAndVarlet/ADASS2019`) to 2026, on Apple Silicon arm64. The original work compared 12 languages and roughly 90 test variants; the present study has grown that to **29 language families and 131 working test variants**, with sixteen language families newly added and a few new variants within existing families. The benchmark itself is unchanged: a tight nested loop that adds the sum of two indices to each element of a 2D floating-point array.

The 2019 paper's sharpest conclusion was the sheer dynamic range of execution speeds — a million-fold spread between fastest and slowest. In the 2026 run the spread is still enormous, about 4.8×10^5 between hand-written NEON[1] assembler and the R reference-class outlier. The qualitative tier ordering is unchanged; what is new is the *populated middle*. Modern compiled languages (D, Zig, Crystal, Nim) cluster in a tight band just below optimised C; JIT-compiled VMs (Java, C#, Go, Kotlin, F#, LuaJIT) form their own near-C cluster; and there is now a distinct "JIT-aided dynamic" tier (PyPy, Ruby +YJIT, V8) that did not exist in 2019. Hand-written arm64 NEON assembler beats every compiler by about 17–19%, a transient state on novel hardware.

1 Background and methodology

The benchmark kernel is unchanged from 2019:

```
for (int Iy = 0; Iy < Ny; Iy++)
  for (int Ix = 0; Ix < Nx; Ix++)
    Out[Iy][Ix] = In[Iy][Ix] + Ix + Iy;
```

Default array shape is 2000 columns × 10 rows = 20 000 elements per call, repeated typically 100 to 5 000 000 times depending on the language (the slowest cases run as few as 3 times) so that the timed run lasts at least several seconds. The harness subtracts a single-iteration run from the timed run to correct for process startup and result-checking overhead.

Hardware is Apple Silicon (M-series), arm64 — *not* the x86-64 used in 2019. This is the most important methodological difference, and matters for two results in particular: the hand-written assembler (was x86 AVX in 2019; now arm64 NEON) and the `-march=native` flag, which in 2019 unlocked AVX-2 and on M-series lands on NEON's 128-bit vector unit instead.

The full set of toolchains is installed locally; the test harness (`Run.py`) sets up paths and environment variables so that all 29 toolchains resolve without further activation. Pre-existing system tools (system clang/g++, gfortran, Perl, etc.) account for about a third; the rest live

*Created PDL in 1996 and may be biassed.

[1]NEON is ARM's SIMD vector extension, the arm64 counterpart of x86's SSE/AVX.

under `tools/` or in dedicated conda environments. The full provenance is documented at the top of `Run.py`.

All numbers in this report come from a single reference run executed on a freshly-rebooted machine on AC power with no other significant activity (`run22.log`).

A note on the "raw", "arrays" and "vectors" labels

Several languages with array libraries (Python+numpy, Perl+PDL, Ruby+Numo, Octave) appear in the benchmark with multiple variants. Following Shortridge (2019), we use three labels for these:

- **"raw"** — per-element access through the array library's machinery, i.e. a nested `for` loop with the host language doing the iteration:

  ```
  out[iy,ix] = ina[iy,ix] + ix + iy                              # numpy
  $Output->set($ix,$iy, $Input->at($ix,$iy) + $ix + $iy)     # PDL
  ```

  This is *not* how these libraries are intended to be used. Each scalar access crosses the language → C boundary, so you pay both the host interpreter cost *and* the array library's slice/box/dispatch cost. Empirically this is *slower* than the pure-language version (Python lists, Perl plain arrays): 644 × for Python lists vs. 2 883 × for numpy raw; 716 × for Perl raw vs. 9 207 × for PDL raw (after the 2026 `->at/->set` correction). The "raw" label is something of a trap.

- **"vectors"** — row-at-a-time slice arithmetic with sequence vectors. Crosses the boundary $O(N_y)$ times rather than $O(N_x N_y)$. Substantially faster than "raw".

- **"arrays"** — a *single* expression that broadcasts two index vectors over the input array:

  ```
  out[:] = ina + arange(nx)[NA,:] + arange(ny)[:,NA]    # numpy
  $Output .= $Input + sequence($Nx) + sequence($Ny)->dummy(0,1)   # PDL
  ```

  Crosses the boundary *once per call* regardless of $N_x N_y$. This is the array-language idiom proper, and is algorithmically identical to the J kernel `y + (i. ny) +/ i. nx`. The three implementations (numpy "arrays", PDL "arrays" and J) compute the same broadcast of two index ranges over the input array.[2]

These labels apply only to the array-library tests. The plain "Python lists" and "Perl raw" tests use the host language's ordinary array structures (Python lists, Perl arrays of arrays) and do not cross any C-extension boundary per element — they are pure host-language interpreter benchmarks.

2 New languages added in 2026

The new entries fall naturally into four groups.

[2]Shortridge (2019) credits Karl Glazebrook for both the numpy and PDL versions of this broadcast trick. A natural alternative in PDL is to use `xvals($Input) + yvals($Input)`, which reads more clearly than `sequence($Nx) + sequence($Ny)->dummy(0,1)` and avoids the slightly opaque `dummy(0,1)` call. We measured this variant directly against the canonical form: it runs about 2.6× *slower*. The reason is allocation cost. `xvals/yvals` each return a fresh *full-shape* 2D piddle ($N_x \times N_y$ elements) materialised before the addition starts, while `sequence + dummy` allocates only two short 1D vectors ($N_x + N_y$ elements) and lets PDL's broadcasting machinery thread the addition over the input. For this benchmark the readability cost is paid in memory bandwidth.

AOT systems languages (5). **D** (LDC `-O3 -release -boundscheck=off`), **Zig** (`-O ReleaseFast`), **Crystal** (`--release`), **Nim** (`-d:danger --gc:arc`), and **OCaml** (`ocamlopt -O3 -unsafe`). All are LLVM-class compiled-to-native[3], sitting in tiers 3–4.

JIT VMs and managed AOT (5). **C# / .NET 8** (JIT and AOT via `PublishAot=true`), **F# / .NET 8** (one variant), **Kotlin** (kotlinc + JVM), **Go** (default and `-gcflags=-B` for no bounds checks), and **Haskell** (GHC `-O2`).

Dynamic languages with non-trivial runtime stories (6). **Ruby** (five variants: raw 2.6, raw 3.3, raw 3.3 +YJIT, Numo, Numo +YJIT), **Lua / LuaJIT** (the tracing-JIT exemplar), **J** (jconsole; the array-language exemplar), **PyPy** (Python lists and numpy raw under PyPy's tracing JIT), **Common Lisp / SBCL** (typed and untyped), and **Octave** (a Matlab-class numerical interpreter, tested with both element-wise loops and vectorised broadcasting — two idioms that turn out to be 800× apart).

New variants within existing families (4). **Javascript**: Bun and Deno (Node.js was already in the suite); **Perl/PDL**: PDL::PP via `Inline::Pdlpp`, plus a separate `->at/->set` variant of the per-element "raw" test (discussed below in §6); **Assembler**: arm64 NEON port replacing the 2019 x86 AVX implementation.

That is sixteen new language families and four new variants within existing families.[4]

3 Updated results — direct successor to Shortridge 2019 Figure 1

Figure 1 is the direct successor to Figure 1 of Shortridge (2019). All 131 test variants are placed on a language × compiler-option heatmap, colour-coded logarithmically from green (fastest) through yellow, orange and red to purple (the R reference-class outlier). The reference value is the hand-written arm64 NEON assembler. Lower numbers are faster.
The visual key:

Tier	Range	Examples
1	1.0 – 1.1	arm64 assembler
2	1.1 – 2.5	Julia int32, fastest C, Fortran `-O3 -march=native`
3	2.5 – 7	Julia, D, Zig, Crystal, Nim, Java, C#, Kotlin, Go, numba, Boost-no-assert, PDL::PP, Rust iter
4	7 – 15	F#, LuaJIT, J, Haskell, OCaml, SBCL typed, Rust safe/unsafe, Boost-with-asserts `-O3`, Bun
5	15 – 100	numpy, Numo, PDL arrays/vectors, Octave vectorised, Swift `-O`, Node, Deno, unoptimised C/Fortran
6	100 – 600	R outer, SBCL untyped, Ruby +YJIT, Rust unoptimised
7	600 – 15 000	**Python lists Py3**, Perl, Ruby raw, R raw, unoptimised Boost, Python lists Py2, Python numpy raw, Swift unoptimised, PDL raw, Tcl
8	≥ 20 000	PyPy + numpy, Octave loop, R ref-class

A per-language-family view of the same data, with min/max ranges and geometric means across each family's variants, is shown in Figure 2. The distribution of all 131 results across the eight tiers is shown in Figure 3, and a log-binned histogram of the 131 individual values in Figure 4.

[3]AOT = Ahead-Of-Time compilation: source → native machine code at build time, before the program runs. Contrast with JIT (Just-In-Time): source → bytecode at build time, then bytecode → native at run time, by the running VM.

[4]Two minor counting choices: "Lua / LuaJIT" and "Common Lisp / SBCL" are each one family; PyPy is treated as a separate family because its interpreter is a complete reimplementation distinct from CPython.

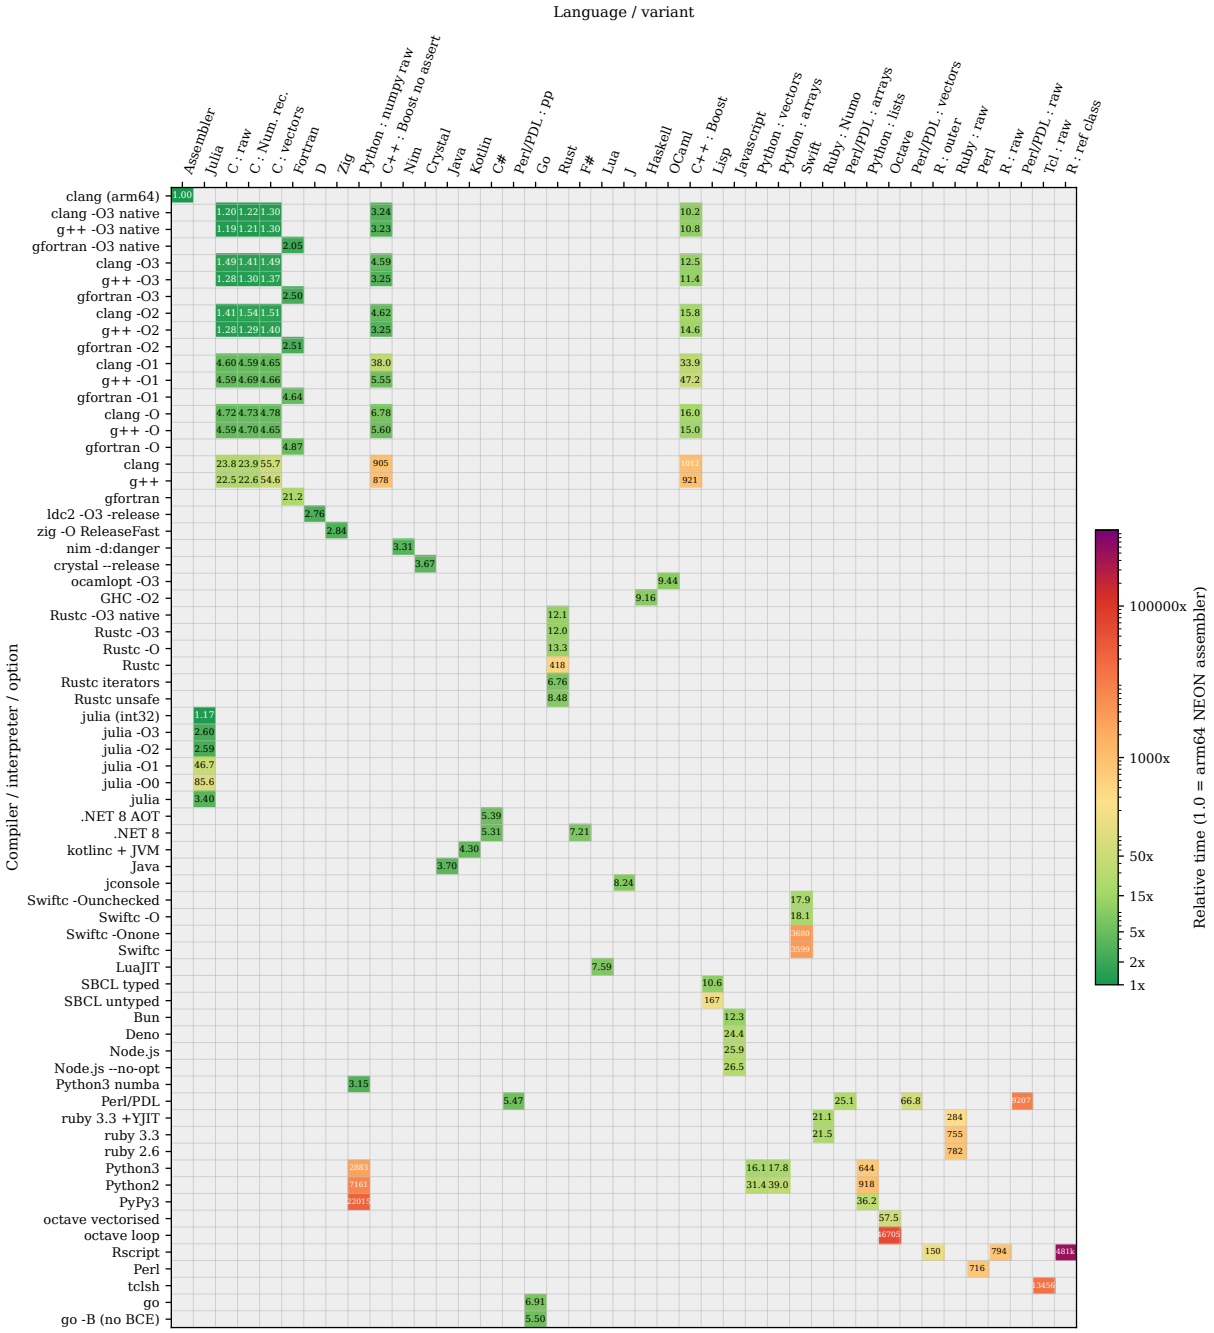


Figure 1: The 2026 successor to Shortridge (2019) Figure 1. All 131 test variants on the same colour scale. Reference: hand-written arm64 NEON assembler. Empty cells are language/compiler combinations not run.

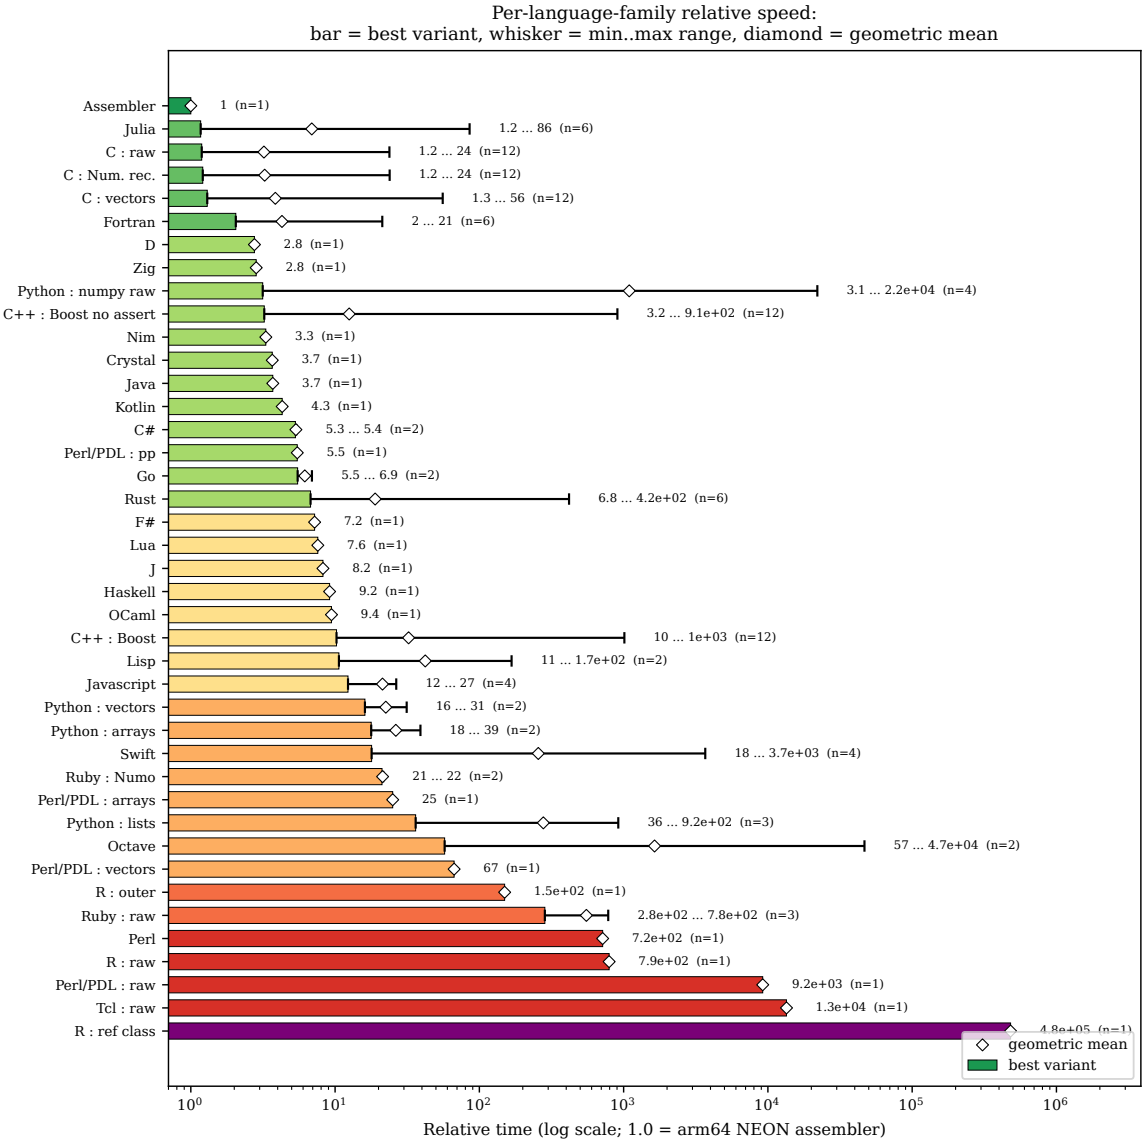


Figure 2: Per-language-family aggregate. The bar shows each family's *best* result; the whisker spans min..max across all variants in that family; the diamond marks the geometric mean. Annotations give the range and the number of variants (n). Multi-variant families (C, Boost, Fortran, Julia, Rust, Python: numpy raw) demonstrate how compiler/optimisation flags can move a single language across two or more orders of magnitude — but the *best* variant always determines the language's tier.

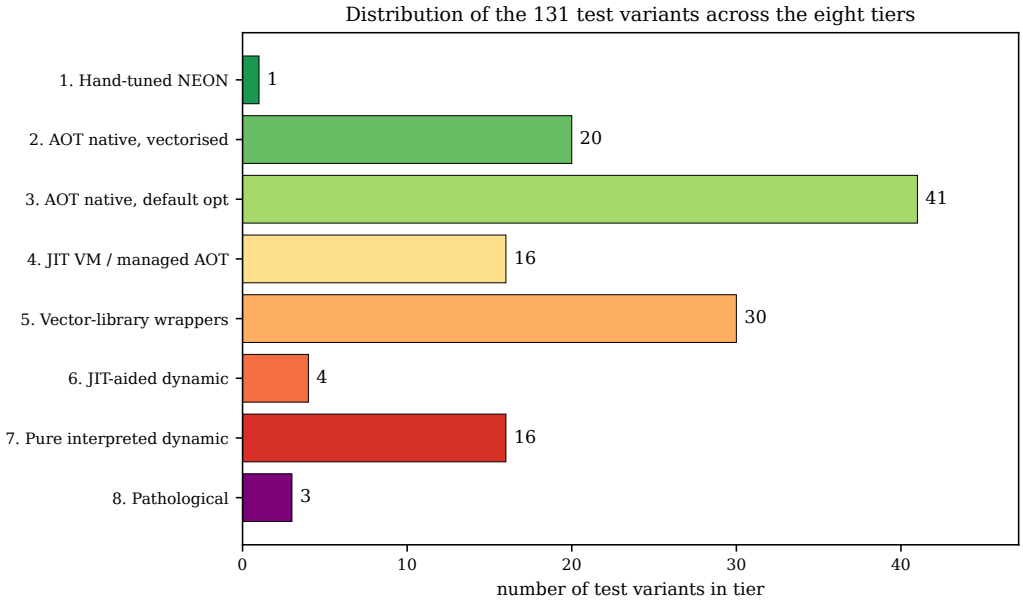


Figure 3: Distribution of the 131 test variants across the eight tiers (one bar per tier; counts annotated). The compiled-to-native tiers (2 and 3) hold 47% of all results; only one entry sits in tier 1 (hand-tuned NEON), and three sit in tier 8 (pathological).

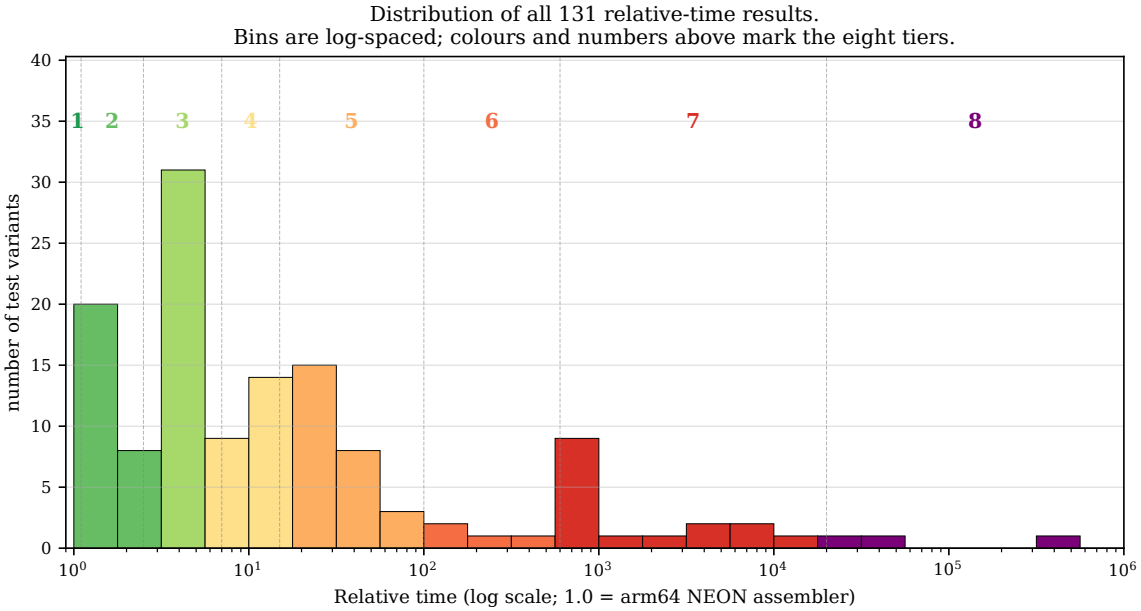


Figure 4: Histogram of the 131 individual relative-time values on a log scale. Bars are coloured by tier; tier numbers are repeated above the histogram for reference. The bimodal shape is clear: a large mass of compiled-to-native results in tiers 2–5 (relative speed 1–100×), and a separate dynamic-interpreter cluster in tier 7 around $10^3\times$. The three pathological tier-8 outliers and the single tier-1 hand-tuned assembler are visible at the extremes.

4 How the languages cluster

The 131 results sort cleanly into the eight tiers shown above. Tier boundaries are not arbitrary; each represents a concrete architectural decision about how the language compiles or interprets numeric loops.

Tier 1 — hand-tuned NEON (1.00). Just the one entry: `cnrsub_arm.s`, written by Shortridge in 2022 as a learning exercise. It beats the next-fastest result (Julia int32 at 1.17×) by about 17%, and the fastest C configuration (1.19×) by about 19%. This is a transient state: M-series clang's auto-vectoriser will close the gap in a release or two. In 2019 hand-written x86 assembler could *not* beat compilers, and that finding will return.

Tier 2 — AOT to native, vector-instruction-using (1.1 – 2.5). Just C, C++, Fortran and Julia, when fully optimised with vector instructions (`-O2`/`-O3` with `-march=native` or equivalent) — plus Julia's int32 trick at the very top. Across this whole tier, the optimised back-end produces almost identical machine code; the language is doing little work the compiler cannot undo. A 2019 finding that survives unchanged: most C/C++ approaches produce essentially the same number once optimisation is turned on. The other systems-tier languages split just below this. D, Zig, Crystal and Nim land in tier 3, alongside the JIT-VM languages (Java, C#, Kotlin, Go). F# at 7.21× has crossed just over the tier-3/4 boundary into tier 4 (in 2019 it was tier 3). Optimised Rust (in its iterator, unsafe and `-O3 native` forms), OCaml and SBCL-typed all sit slightly lower in tier 4, because bounds checks, array layout or type-information overhead leaves a small but visible cost on this kernel. SBCL-untyped is a different case entirely and falls all the way to tier 6.

Tier 3 — compiled or JIT-compiled to native, no vector hints (2.5 – 7). Most of the dataset's compiled languages plus the JIT VMs land here: D, Zig, Crystal, Nim, Java, C# (JIT and AOT — they coincide here), Kotlin, Go, plus C/Fortran at `-O`/`-O1` and the various Boost-no-assert configurations. This is the most populous tier of the suite (41 entries). What unites it: the inner loop is compiled to native code at some point, but without the vector-instruction hints that move tier 2 ahead. D (2.76) and Zig (2.84) sit at the top, with numba (3.15), Nim (3.31), Crystal (3.67) and Java (3.70) close behind; Go default (6.91) is near the bottom. F# at 7.21× has crossed just over the boundary into tier 4. Notably, **Perl/PDL::PP (5.47) lives in this tier** despite being a Perl test — because its kernel is compiled C lowered through PDL::PP and called from Perl as a single dispatch, the Perl-the-language wrapping costs are amortised away (see §6 for the syntax detail).

Tier 4 — harder cases of JIT and AOT (7 – 15). LuaJIT (7.59), J (8.24), Rust safe and unsafe at `-O3 native`, Haskell, OCaml, SBCL-typed, Bun, and the C++ Boost variants *with* assertions on. What unites this tier is the presence of *some* runtime check or abstraction the optimiser cannot fully elide: Rust's bounds-check elision is incomplete in the indexed-loop form; Boost's assertions add per-element checks; J's primitive dispatch has unavoidable per-call overhead; OCaml's `float array array` is row-pointer-indirect rather than contiguous. All of these compile to native code, but none of them reach the unconditional fast path of tier 3. (PDL::PP at 5.47 sits just below this tier — close enough that the boundary feels slightly arbitrary; the tier 4 entries are paying ∼2× of optimal-C overhead from runtime checks, while PDL::PP is paying ∼1× from its broadcasting macros.)

Tier 5 — Vector-library wrappers (15 – 100). numpy "vectors"/"arrays", Ruby Numo, Perl/PDL arrays/vectors, Octave vectorised. The kernels are compiled C; what is measured is the per-element overhead of the vector library's broadcasting machinery (PDL_Indx arithmetic, slice descriptors, dtype dispatch, the GIL on Python's side). Same floor everywhere — the host language is not the limiting factor.

Tier 6 — JIT-aided dynamic (100 – 600). Just four entries: R outer (150), SBCL untyped (167), Ruby +YJIT (284), and unoptimised Rust (418). A small mixed tier of partially-compiled-but-not-fully or compiled-but-not-optimised configurations. Ruby +YJIT is the prototypical entry: a method-level JIT that gets ∼2.7× over plain Ruby but lacks the analysis depth of LuaJIT or V8.

Tier 7 — Pure interpreted dynamic (600 – 15 000). Sixteen entries, of which the most representative are: **Python lists CPython 3 (644)**, Perl (716), Ruby raw 3.3 (755), R raw (794), unoptimised Boost C++ (878–1012), Python numpy raw Py3 (2 883), Swift unoptimised (∼3 600), **PDL raw (9 207, after the 2026 `->at/->set` correction)**, and Tcl (13 456). Bytecode dispatch per operation, dynamic dispatch per arithmetic op, no specialisation across types. The ∼20× spread within this tier reflects how heavy the per-bytecode-op work is, not anything fundamental about the languages' type systems. Bytecode dispatch per operation, dynamic dispatch per arithmetic op, no specialisation across types. The 20× spread within this tier reflects how heavy the per-bytecode-op work is, not anything fundamental about the languages' type systems.

Tier 8 — Pathological (≥ 20 000). Octave loop, R ref-class, PyPy + numpy. Each of these is a *legitimate* coding choice that interacts catastrophically with the language runtime's design. They are the most pedagogically useful results in the suite — if you understand *why* each is in this tier, you understand the language much better.

5 The vectorised-language picture

Four language families (Python, Perl, Ruby, Octave) have one or more "vector-library" idioms alongside their conventional one, giving six rows of comparison in Table 1. In each case there is a dramatic difference.

Language	Loop tier	Vector tier	Speed-up
Python (numpy "vectors")	644 (lists)	16	40×
Python (numpy "arrays")	644 (lists)	18	36×
Perl (PDL arrays)	716 (pure Perl)	25	29×
Perl (PDL vectors)	716 (pure Perl)	67	11×
Ruby (Numo)	755 (pure Ruby)	22	35×
Octave	**46 705 (loop)**	**57 (vec)**	**813×**

Table 1: Intra-language speed-ups from switching to the vectorised idiom. Octave is in a class of its own.

Octave's 813× spread between the loop and vectorised forms is the dramatic centrepiece of the table (Figure 5). Both forms are *idiomatic Octave*: a working scientist would naturally write either one. Yet the loop form is among the slowest results in the entire 131-test suite, while the vectorised form sits in the same tier as numpy and Numo. (Perl/PDL has a smaller but still wide same-language gap — the per-element PDL "raw" form (9 207×, after the 2026 `->at/->set` correction) versus the inline-C PDL::PP kernel (5.47×) spans about 1 700× — but the two forms there are not "loop" vs "vectorised" in the same idiomatic sense; PDL::PP is a compiled C kernel embedded in a Perl script.)

A separate question is how the vector-library kernels compare to each other. Once you cross into the compiled C extension, the host language stops mattering:

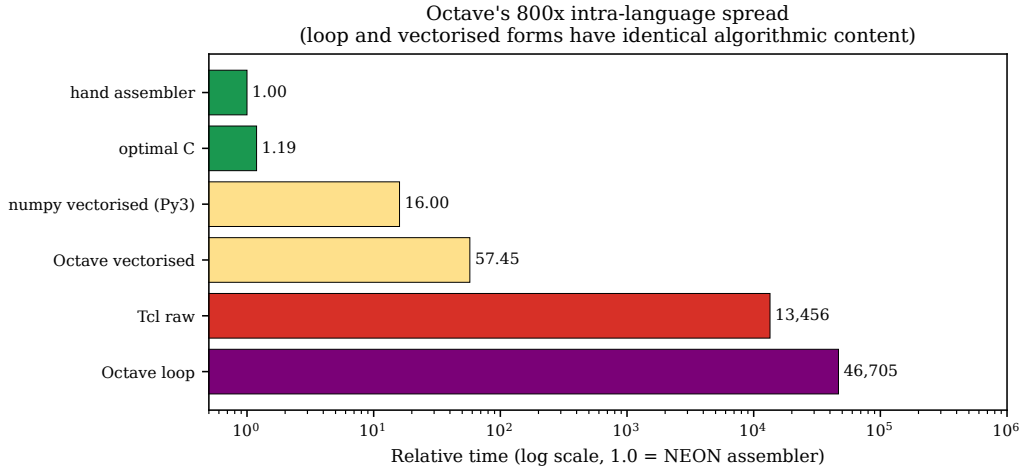


Figure 5: Octave's ∼810× intra-language spread. Both implementations are idiomatic Octave; both compute the same answer. The difference is whether the work is done by Octave's interpreter (46 705×) or its compiled vector kernel (57×).

Library	Host	Rel. speed
PDL::PP (compiled C kernel)	Perl	**5.47**
numpy "vectors"	Python	16.0
numpy "arrays"	Python	17.6
Ruby Numo	Ruby	21.5
Perl/PDL arrays	Perl	25.1
Octave vectorised	Octave	57.5
Perl/PDL vectors	Perl	66.9

The ordering is determined almost entirely by the cleverness of the library, not the host language. PDL::PP sits ahead of every numpy result, every Numo result, and every JavaScript engine. The Perl-the-language overhead is not the limit; PDL's broadcasting machinery is.

A particularly interesting case is the **J language** (tier 4, 8.24×). J is a successor to APL: dynamically typed, interpreted, with a heavily symbolic syntax — yet it lands ahead of Bun (12.29×), Node.js (25.88×), and Ruby Numo (21.54×, or 21.15× with YJIT). The whole benchmark is one expression:

```
out =. y + (i. ny) +/ i. nx
```

("outer-product-add the row indices and column indices to make an $iy + ix$ matrix, and add it to the input array y"). Because the entire 20 000-element kernel runs as one primitive call, J's interpreter dispatch overhead is amortised away. J pays the same "cross-into-compiled-C-once" cost as numpy and Numo — but its broadcasting machinery is slimmer than the high-level array-library wrappers, putting it ahead of every numpy/Numo/PDL-array result in the suite. (J does not beat PDL::PP, which lands at 5.47×: PDL::PP is no longer wrapping at the array level but compiling a small C kernel inline.)

The lesson, consistent with 2019's observation about numpy: *the vector-library tier is governed by the library, not the language.* Choose your scientific stack on language ergonomics, and trust that the underlying tensor library will produce roughly the same number.

6 Notable syntaxes and language-specific optimisations

J — array-orientation taken to the limit. The full J port is nine lines (most of which are argument parsing and result-checking). The kernel itself is one expression: `y + (i. ny) +/ i. nx`, where `i.` is range, `+/` is reduce-sum, and adjacent verbs compose. J's roots are in Iverson's APL;

the syntax is dense to the point of hostility but the *thinking* it encourages — operate on whole arrays, never write a loop — is modern data science distilled. The performance result (8.24×) is largely a side-effect of this discipline: when there are no loops, there is nothing for the interpreter to slow down.

SBCL Common Lisp — typed vs. untyped (15× difference). Two source files, identical except the typed variant adds:

```
(deftype matrix () '(simple-array double-float (* *)))
(declaim (ftype (function (matrix fixnum fixnum matrix) (values))
                subr))
(defun subr (in-arr nx ny out-arr)
  (declare (type matrix in-arr out-arr) (type fixnum nx ny))
  ...)
```

Untyped: 167.25× relative time (tier 6, alongside Ruby +YJIT and R outer). Typed: 10.63× (tier 4, alongside Haskell, OCaml and LuaJIT). This is the most pedagogical single result in the dataset. SBCL is a sophisticated whole-program compiler — but it cannot know what type a variable is unless told. With a few `(declare (type ...))` forms, generic arithmetic dispatch becomes inlined double-float arithmetic and the inner loop becomes tight machine code. Without them, every `+` is a runtime polymorphic dispatch through SBCL's generic-function tables.

A **16× speed-up for adding eight declarations to a working program**. This is exactly the value proposition Julia inherits from Lisp: write fully untyped code and it works; add types to the hot loop and it gets fast.

PDL::PP — compiled C inside a `.pl` script. PDL::PP compiles a small C kernel via `Inline::Pdlpp`. The kernel itself is:

```
pp_def('csub',
   Pars => 'i(m,n);[o] o(m,n)',
   Code => '
      loop(n) %{
          loop(m) %{
            $o() = $i() + m+n;
          %}
      %}
   ');
```

The Perl wrapping is essentially free — once compiled, calling `$InArray->csub` dispatches into the compiled `.bundle`. The 5.47× result is determined entirely by PDL's broadcasting macros: `loop()` expands to a PDL_Indx-aware iterator with slice-descriptor checks per element, not a flat C `for`.

PDL "raw" — a 2019 implementation choice corrected in 2026. The original `csubpdl.pl` in Shortridge (2019) implemented per-element access using `$Output->slice($ix,$iy) .= ...`. That choice was a mistake. `->slice` constructs a one-element *slice piddle* on every read *and* every write — the slice machinery is designed for sub-array views, not for scalar access — which dominates the per-element cost and produced an artificially bad result of $\sim 10^5\times$ relative time, by far the worst Perl/PDL number in the 2019 figure. The *intended* primitive for scalar access in PDL is `->at($ix,$iy)` (read, returns a Perl scalar) and `->set($ix,$iy,$v)` (write, no slice piddle constructed). This is the proper PDL analogue of numpy's `ina[iy,ix]` indexing and is what Shortridge meant the test to do.

This 2026 reprise replaces the 2019 form with the `->at`/`->set` form. The corrected PDL raw result is **9 207**×, about 11× faster than the 2019 `->slice` version, and now in the same broad ballpark as Python numpy raw (2 883×, on CPython). It is still slow — you are still looping in

Perl and crossing into C on every element — but it is no longer worst-of-suite. PDL raw now sits in tier 7 alongside Tcl, R raw, Python lists Py2, and unoptimised Boost; tier 8 (pathological) now contains just three entries — PyPy+numpy (22 015), Octave loop (46 705) and R reference-class (481 009). The 2019 paper's "the per-element PDL access form is the second-worst result in the suite" is no longer true under this corrected implementation.

The original `->slice`-based test has been removed from the benchmark suite; `csubpdl.pl` now contains the `->at`/`->set` form. (The `csubpdl.pl` source file in the repository carries an explanatory comment at the top documenting the change.)

numba — the Python decorator that JITs.

```
from numba import njit

@njit
def subr(In, Nx, Ny, Out):
    for iy in range(Ny):
        for ix in range(Nx):
            Out[iy, ix] = In[iy, ix] + ix + iy
```

One decorator. 2 883× → 3.15×: a 915× speed-up by adding a single decorator to a numpy-using function. **Note that numba is not really doing anything *numpy*-specific here**: it is JIT-compiling a regular Python loop, treating the numpy array as a typed contiguous double buffer. It does not call into numpy's compiled C kernels at all. The "numpy" in the test name simply provides a strongly-typed contiguous storage that numba's type-inference can specialise for; the loop itself is then lowered to LLVM IR and from there to native machine code that completely bypasses numpy's Python boundary. The result is that numba performs essentially the same job that LuaJIT or a Crystal compiler would perform on equivalent code: it turns Python's per-element interpreter dispatch into a flat native loop. This is the single most cost-effective speed improvement available to a working scientist who has Python code in hand and a five-minute time budget.

Crystal — Ruby syntax with types. Identical-looking source code to Ruby, but Crystal infers types at compile time and lowers to LLVM IR. Result: 3.67× — *near C-class*. The same algorithm in Ruby (with YJIT) is 284×; without YJIT, 755× (Ruby 3.3) or 782× (Ruby 2.6). Crystal demonstrates that **the cost of being Ruby-the-language is not the syntax; it is the dynamic dispatch**.

Kotlin = Java, F# ≠ C# (but close). Kotlin (4.30) lands slightly slower than Java (3.70) on the same JVM. F# (7.21) lands ∼36% slower than C# (5.31) on the same .NET runtime. Both are explainable: Kotlin's nullability-checked types compile to a few additional method calls the JIT can inline *almost-but-not-quite* away; F#'s discriminated unions and option types add small constants to imperative loops. **Within runtime-pairs, the source language matters at the 20–40% level, the runtime matters at the order-of-magnitude level.**

LuaJIT and J — the frontier of "dynamic but fast". LuaJIT (7.59) and J (8.24) are the two cleanest demonstrations that a dynamically-typed language can run at AOT-compiled speeds. They get there by entirely different means:

- **LuaJIT** is a *tracing JIT*: it observes hot loops, infers type-specialised machine code for them, and emits inline assembly. The user writes ordinary Lua; the JIT does everything.
- **J** is *not* JIT-compiled at all — it is interpreted, but its primitive operations are themselves tight C kernels and the *language* discourages writing per-element loops. The user gets speed by accepting the array-oriented idiom.

These are the two primary paths to closing the dynamic-language performance gap. PyPy (Python's tracing JIT) takes the LuaJIT path but lands further behind, because Python's dynamic semantics are considerably more complex than Lua's, leaving the JIT less to work with.

7 Anomalies

The benchmark surfaced six results that were not predictable from first principles, all of which are pedagogically valuable.

PyPy + numpy is 600× *slower* than PyPy + Python lists.

Configuration	Rel. speed
PyPy + Python lists (`csubraw.py`)	36
CPython + numpy raw (`csubnp.py`)	2 883
PyPy + numpy raw (`csubnp.py`)	**22 015**

This is the single largest counter-intuitive finding in the new data. Almost every Python programmer has heard "PyPy is faster than CPython"; almost none know that this is *false* for numpy-heavy code, sometimes by an order of magnitude. The mechanism: PyPy implements its own internal object model, designed for the tracing JIT to reason about. To support compiled extensions like numpy that target the CPython C API, PyPy provides a shim called `cpyext` that translates between PyPy objects and the shapes numpy expects. Every call into a numpy C extension crosses `cpyext`, which is much slower than CPython's native fast path. For tight loops that invoke numpy thousands of times, the per-call `cpyext` cost dominates. The JIT cannot help because the JIT cannot trace into compiled C.

Working scientists' takeaway: **PyPy is only beneficial for *pure* Python code**. As soon as you cross into a numpy/scipy/scikit C extension, it becomes a net slowdown. CPython remains the right choice for the numpy-using majority of scientific Python.

Octave loop — worst-of-suite performance from idiomatic code. Octave's element-wise loop runs at 46 705× relative time — slower than Tcl, slower than R-raw, slower than PyPy+numpy (22 015×), only beaten in the bottom of the suite by R ref-class (481 009×). Yet the source is *idiomatic*; see Figure 5. The hot kernel of `csubraw.m` is the obvious nested-loop translation:

```
for iy = 0:ny-1
   for ix = 0:nx-1
      out(iy+1, ix+1) = ina(iy+1, ix+1) + ix + iy;
   end
end
```

The vectorised form (`csubvec.m`) does the same thing in one broadcast. Outside the timed loop, two short index vectors are built once — a $1 \times N_x$ row and an $N_y \times 1$ column:

```
ix_row = 0:nx-1;       % 1 x nx (row vector)
iy_col = (0:ny-1)';    % ny x 1 (column vector, transposed)
```

The hot kernel is then a single expression. Octave broadcasts the row across all rows of `in`, the column across all columns, and adds:

```
out = in + ix_row + iy_col;
```

That single expression runs 813× faster (57× relative time) than the nested loop above. **Octave is a domain-specific language for vectorised arithmetic**; the scalar interpreter exists because it must, not because it should be used. Modern Matlab (R2015b+) ships a JIT that specifically targets this anti-pattern — Octave does not.

Perl/PDL raw and Python numpy raw — traps disguised as idiomatic code. Perl/PDL raw lands at 9 207× (after the 2026 `->at/->set` correction), and Python numpy raw at 2 883× — markedly worse than their *plain* pure-language counterparts (Perl raw 716×, Python lists 644×). The Perl factor is the more dramatic of the two: PDL raw is ∼13× worse than pure Perl, while CPython numpy raw is ∼4.5× worse than CPython lists. Both are surprising results until one notices what "raw" actually means in those tests: **a host-language loop that calls into the C array library on every single element**. For PDL, each iteration calls `$ina->at($ix,$iy)`, which crosses the Perl-to-C boundary, dereferences the piddle's storage, and returns a Perl scalar — and then `$out->set($ix,$iy,$v)` does the inverse on the write side. For numpy, each `ina[iy,ix]` traverses numpy's `__getitem__` machinery, builds a Python scalar from the underlying C buffer, and returns it through the C-API; the assignment side does the inverse. The host-language interpreter overhead is then *added* to that, not amortised away. Both "raw" tests are slower than the equivalent code written using only the host language's native data structures, because they pay the host's per-iteration cost *plus* the C-extension's per-call cost. The "raw" label is misleading: these are not vectorised operations in the J/PDL-arrays/numpy-arrays sense (one C-side broadcast); they are scalar accesses into a vectorised data structure, which is the worst of both worlds. The only thing the user is saving by using PDL/numpy in this style is the in-memory layout (a contiguous C buffer) — and on this benchmark, that saving is not enough to pay for the per-element traversal overhead.

R reference class — 481 009× (the 2019 outlier). Inherited from the 2019 study with the same root cause. R's reference classes are closure-based; modifying a single field copies the entire encapsulated state. For an array-modifying method called per-element, this means an $N \times N$ copy per element write — quadratic-times-iterations work for what should have been constant-per-call work. The 2019 paper reported about 1.5 M× on x86; the 2026 number is about 481 000× on arm64 (a faster assembler baseline divides into a different relative number). The lesson is unchanged: *idiomatic code can collide catastrophically with a language's runtime model.*

OCaml — `Bigarray.Array2` slowdown vs. `float array array`. OCaml is widely advertised as "within 2–3× of C". The first version of the OCaml port used `Bigarray.Array2` (intended specifically for numerical work) and ran at ∼120× optimal C — a catastrophic gap. Switching to OCaml's native `float array array` (unboxed per-row arrays of doubles) cut the time by ∼13×, landing at 9.44×. The mechanism is per-row allocation overhead: `Array2.slice_left` allocates a fresh `Array1` view per call, defeating tight-loop optimisation. Direct 2D access via `arr.{iy, ix}` is no faster because `Bigarray`'s index-conversion macro is heavier than hand-rolled multiplication. The lesson: even within "fast OCaml," data-layout choice produces order-of-magnitude differences. The documented "near-C" performance in the OCaml literature assumes `float array` not `Bigarray`.

Julia int32 — 1.17× through one type-coercion change. The original Julia test (using floats throughout) ran at about 3.40×. Changing one expression to use int32 arithmetic for the indices — a hint Shortridge added in 2019 in response to ADASS feedback — produces 1.17×. **A two-line code change reaching the second-fastest result in the entire suite, beaten only by hand-written assembler.** This result most strongly supports the 2019 paper's recommendation of Julia for new scientific work.

D and Zig — modern systems languages, identical performance. D's LDC compiler and Zig's stage 0 backend are both LLVM-based, both producing machine code that lands within ∼3% of each other (D 2.76×, Zig 2.84×) — within the measurement noise of the same number. Two languages with very different design philosophies (D: optional GC, ranges, template-heavy; Zig: manual memory, comptime, no hidden control flow) compile to essentially the same SSA-IR

shape on this benchmark after the LLVM optimiser runs. The implication: **LLVM-IR-emitting languages are converging at the back-end**. Differences in front-end design rarely survive `-O3`.

Swift — a ∼200× optimisation cliff. Swift's relative numbers are striking: `swiftc` (default, which is `-Onone`) lands at 3 599×, `-Onone` explicit at 3 680×, but `-O` brings this down to 18.14× and `-Ounchecked` to 17.91×. That is a ∼200× swing from one optimisation flag — one of the largest `-O0`/`-O` gaps of any language in the suite, comparable to C++ Boost-no-assert (which goes from ∼900× at `-O0` to ∼3× at `-O3 native`, a ratio of ∼280×). The reason is structural: Swift's runtime carries several costs that are normally invisible because the optimiser eliminates them, but which dominate when it does not. An idiomatic `Array<Array<Float>>` element access at `-Onone` pays:

- **ARC retain/release** on the inner row reference each time the subscript is taken (Swift's automatic reference counting is a runtime call, not free);
- **Generic dispatch** through the `Collection` protocol witness table on every `subscript(_:)` call — Swift's generics are not specialised at `-Onone`;
- **Bounds checking** on each access;
- **Copy-on-write checks** on the array storage.

At `-O`, the optimiser specialises `Array<Float>.subscript` for the concrete type, devirtualises the Collection-protocol calls, hoists ARC operations out of the loop or elides them entirely, and inlines the inner loop down to plain pointer arithmetic. The 200× spread is the cost of those four runtime mechanisms, end-to-end, with no compiler help. Shortridge (2019) saw essentially the same ∼470× behaviour on the original x86 hardware — this is a property of the language's runtime model, not the platform. The practical takeaway: *never benchmark Swift code without* `-O`; the result will mislead you by two orders of magnitude.

8 Comparison with 2019

A handful of languages were in both studies. Where comparable, the relative-speed ordering is preserved, but several quantitative numbers have shifted. Figure 6 shows the comparison graphically.

The qualitative picture is unchanged. Language-tier ordering is robust across hardware platforms and seven years of compiler evolution. **Where languages have improved, they have done so within their tier**, not by jumping to a higher one. The exception is Python, where numpy's fast path has gotten substantially faster, and (more importantly) numba has emerged as a near-C JIT for numpy code.

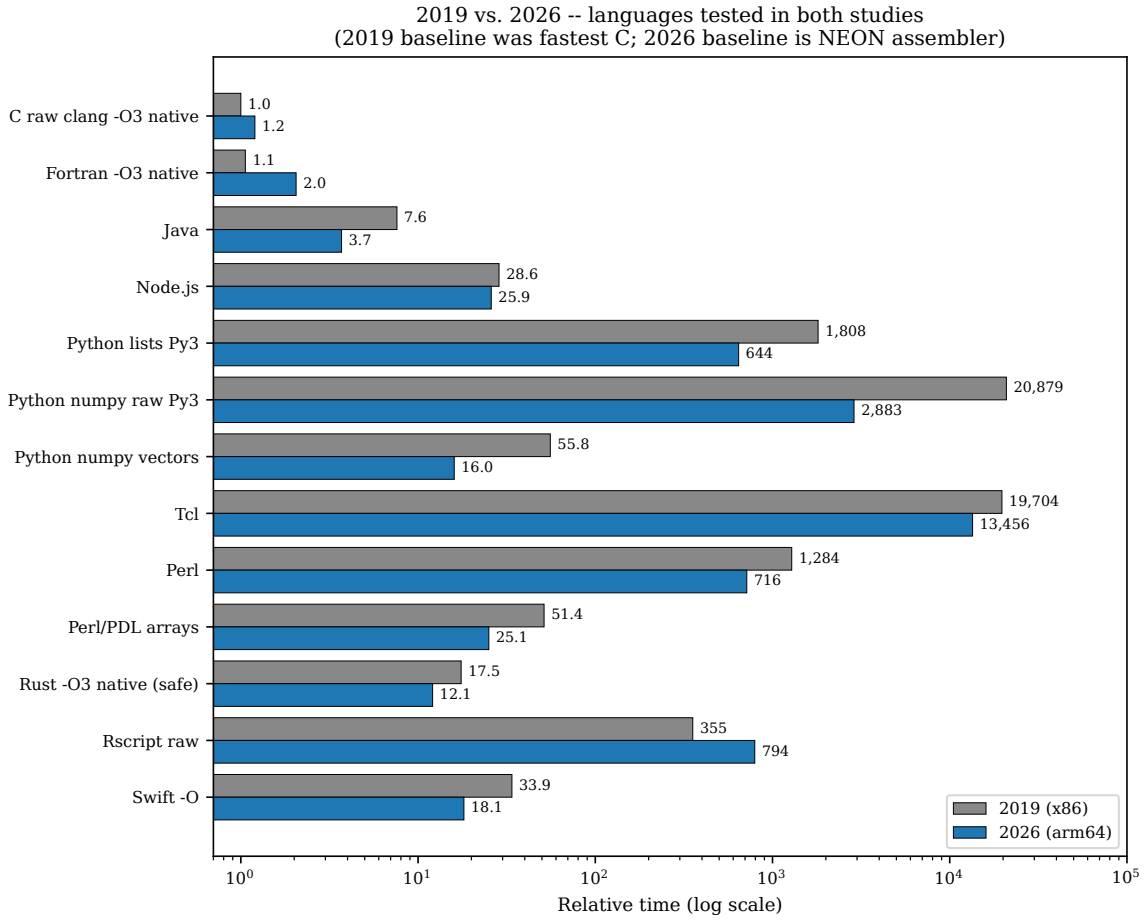


Figure 6: The 13 results that have a direct counterpart in both studies (Shortridge 2019, x86; this 2026 reprise, arm64). The 2019 baseline was the fastest C; the 2026 baseline is hand-written NEON assembler. Most languages stayed in their tier or improved; the qualitative ordering is unchanged.

9 Recommendations by use case

The 2019 paper closed with the modest recommendation to "look closely at the performance you are getting" and to consider Julia or Rust for new scientific work. Seven years on, the dataset supports more specific advice.

Writing the inner loop of an astronomical pipeline. Use **C, Fortran, or Julia** with full optimisation (`-O3 -march=native`). All three sit in tier 2 ($1.1-2.5\times$). The choice between them is ergonomic, not performance-driven. C if you need to interface with FITS libraries or call from Python; Fortran if your group is already invested; Julia if you want an interactive REPL and modern syntax with the same back-end performance.

If you need every last percent of speed and have NEON-shaped data, hand-written arm64 assembly is currently 17–19% faster than the best compilers. This will not last. It is also not a recommendation for production code.

Writing tools / utilities / glue around scientific libraries. Use **Python with numpy** (or PDL if you are already in Perl). Once your hot loop is in a vector-library call, the host language matters very little (tier 5, $15-100\times$). For the small fraction of work where Python's per-element path is hit anyway, add `@njit` from numba — nearly $10^3\times$ speed-up for one decorator, no recompilation, no rewriting.

Do **not** use PyPy if your code calls into numpy. PyPy is a substantial *slowdown* for numpy-heavy work (tier 8). PyPy is only appropriate if your hot loop is genuinely pure-Python.

Prototyping in an interactive scientific REPL. **Julia** (1.1 – 3×) is the best choice the dataset offers. It has the responsiveness of Python, the performance of C, and a single syntax across all tiers (no need to learn numpy idioms or write a C extension). The 2019 paper's enthusiasm has been validated.

For groups already invested in Python, the `numpy + numba + matplotlib` stack remains pragmatic.

Idiomatic OO code with serious performance. **C# / .NET** (5.31×, JIT) or **Java / OpenJDK** (3.70×) are both strong choices. F# (7.21) and Kotlin (4.30) sit just behind their C-family siblings on the same runtimes — choose for syntax, not speed. If your performance budget extends a bit further, **Crystal** (3.67) or **Go** (5.5 – 6.9) are AOT-compiled alternatives.

Modern systems language with memory safety. **Rust** (6.8 – 13.3, depending on idiom) is the obvious choice and has the strongest ecosystem of the new systems languages. **D** (2.76) and **Zig** (2.84) are both *faster* than safe Rust on this benchmark, and demonstrate that a memory-safe language can still land at C speed when the design lets the optimiser see through the abstractions. Rust is the right answer for most groups in 2026; Zig and D are right answers for specific projects.

Teaching what "high performance" looks like. Run your students this benchmark on three Octave configurations: the loop variant, the vectorised variant, and the same algorithm in C. The 813× spread between Octave's two idioms is the most vivid single-language demonstration of "use the language fast path" in any language we tested. Then run it again in SBCL with and without type declarations. The 16× spread there shows the same lesson, applied to type information.

Demonstrating that language design matters. Compare these four results from the dataset:

Python lists, CPython	644
Python lists, PyPy	36
Python lists with `@njit` numba	≈ 3
Hand-written C	1.19

Nearly three orders of magnitude (644 to 1.19 spans a factor of ∼540) for what a working scientist can plausibly think of as "the same problem" in Python, depending on which runtime processes the source. The first three rows are even *the same source file*: switching from CPython to PyPy is zero code change; adding a numba decorator is one line and requires switching from a list to a numpy array. The C version requires a rewrite. This is the broadest single point about "how fast can a language go" the benchmark can demonstrate.

Combinations to avoid. A few combinations the dataset shows are genuinely poor choices for inner-loop scientific work, and which are not just slow but which *surprise* their users with how slow they are:

- Octave element-wise loops (46 705×). Vectorise instead.
- R reference classes for in-place array updates (481 009×).
- Per-element PDL access (9 207×, after the 2026 `->at`/`->set` correction). Use PDL's vectorised forms.
- numpy raw element-wise loops on PyPy (22 015×). Use CPython.

These are pathological *not* because they cannot be made faster but because their idiomatic surface form looks like normal code.

10 Closing thoughts

The 2019 paper's headline — *this code is over 1 000 000 times faster than that code* — was true then. In this 2026 run the corresponding fastest-to-slowest spread is about 481 000×: smaller than the 2019 headline, but still pedagogically the same result. The seven-year gap has widened the variety of points along that range.

What the new languages reveal is more nuanced. The compiled-systems tier has converged: Crystal, D, Zig, Nim, Go, OCaml, SBCL-typed, optimised Rust idioms, and Haskell all sit in a tight band between 2.8 and about 13. Different front-ends, different paradigms (lazy, OO, functional, imperative), but the same LLVM-class back-ends produce the same shape of machine code. Front-end ergonomics now drive the choice, not measured performance.

The dynamic-language tier has bifurcated. CPython, plain Ruby, plain R, Perl, and Tcl remain in the 600 – 15 000× range — slow but at least predictable. Above them, a new "JIT-aided dynamic" zone now spans tiers 4–6: at tier 4, LuaJIT (7.59) and Bun (12.29); at tier 5, PyPy lists (36) and the V8-based JavaScript engines Node (26) and Deno (24); at tier 6, Ruby +YJIT (284) and R outer (150). The most populous tier in the 2026 dataset is tier 3 (41 entries) — the tight band of "compiled or JIT-compiled to native, no vector hints" results into which essentially all of D, Zig, Crystal, Nim, Java, C#, Kotlin, Go, plus C/C++/Fortran at `-O`/`-O1`, now fall. JIT compilation has gone mainstream, and most languages now have *some* path to it.

The vector-library tier (numpy, PDL, Numo, Octave) is unchanged in character but has more entrants. It remains the single highest-value option for scientific Python users — and PDL::PP, in particular, is the highest-leverage way to write a hot kernel inside a `.pl` script.

And, finally, the most pedagogically rich finding of the whole exercise: **language and runtime are separable**. The same language can move by factors of several to nearly three orders of magnitude depending on which runtime, compiler, type information or fast path is used: Ruby raw at 755× vs. Ruby +YJIT at 284× (a factor of ∼2.7); SBCL untyped at 167× vs. SBCL typed at 11× (∼16×); CPython lists at 644× vs. PyPy lists at 36× (∼18×); and CPython numpy raw at 2 883× vs. numba at 3.15× (∼ 915×, the largest such factor in the Python family). Even this is exceeded by the Perl family: PDL raw vs. PDL::PP, after the 2026 `->at`/`->set` correction to the raw form, is 9 207× vs. 5.47× = ∼1 683×, the largest single-language fast-path factor in the dataset. The benchmark makes explicit what scientific programmers often do implicitly when they reach for `numpy`, PDL::PP, numba, type declarations or a JIT runtime.

The original paper's appendix on R reference classes — "a cautionary tale, some attempts to optimise can backfire" — remains the appropriate cautionary tale.

References

- Shortridge, K. 2019. "Compilers, languages, options, and astronomical images." *Astronomical Data Analysis Software and Systems XXIX*, ASP Conference Series, Paper P9-17.
  Conference paper PDF (4 pages): `https://github.com/KnaveAndVarlet/ADASS2019/blob/master/P9-17.pdf`
  Poster: `https://github.com/KnaveAndVarlet/ADASS2019/blob/master/LanguagesAndImageProcessing.pdf`
  Full write-up (53 pages, the per-language deep dive): `https://github.com/KnaveAndVarlet/ADASS2019/blob/master/FullWriteUp.pdf`

- Shortridge, K. 2019. `KnaveAndVarlet/ADASS2019` (GitHub repository). All source code for the benchmark, including the original 2019 language ports, the test harness `Run.py`, and the colour-coded

results figure used as Figure 1 of the conference paper.
`https://github.com/KnaveAndVarlet/ADASS2019`

- Shortridge, K. 2022. `cnrsub_arm.s`. Hand-written arm64 NEON port of the benchmark kernel, contributed to the repository in 2022 as a follow-up to the 2019 study. Used as the 1.00× reference in this 2026 reprise.
- Glazebrook, K. & Soeller, C. 1996–present. **PDL — The Perl Data Language.** First publicly released as PDL 1.00 in October 1996. Project home page: `https://pdl.perl.org/`. "Day 24: Perl Data Language reflections," PDL Advent Calendar 2024: `https://pdl.perl.org/advent/blog/2024/12/24/reflections/`. For background on the language and its origins, see *Perl Data Language* on Wikipedia (`https://en.wikipedia.org/wiki/Perl_Data_Language`).
- Numo (Ruby's numerical array library) — the Ruby analogue of numpy, used for the "Ruby : Numo" tests: `https://github.com/ruby-numo/numo-narray`.

A Files

The complete updated benchmark suite lives in `ADASS2019/` (alongside `tools/`, `logs/` and `reports/` as siblings of that directory):

- `Run.py` — the harness, with bootstrap for `tools/`-installed toolchains and conda environments.
- `cnrsub_arm.s` — Shortridge's hand-written arm64 NEON port (the reference implementation).
- Source files: `csubraw.cs`, `csubraw.fs`, `csubraw.go`, `csubraw.zig`, `csubraw.d`, `csubraw.nim`, `csubraw.lua`, `csubraw.kt`, `csubraw.rb`, `csubraw.ml`, `csubraw.lisp`, `csubraw-typed.lisp`, `csubraw.hs`, `csubraw.cr`, `csubraw.m`, `csubvec.m`, `csubpdlpp.pl`, `csubnumo.rb`.
- `run22.log` — the canonical reference run (this report's source numbers).
- `tools/` — local installs (Julia, Rust, Bun, Deno, Go, LuaJIT, .NET 8, Crystal, J, PyPy, Zig, LDC, Nim, GHC, custom YJIT-built Ruby, Boost headers).
- Conda envs: `adass-r`, `adass-py2`, `adass-rb`, `adass-sbcl`, `adass-ocaml`, `adass-kotlin`.
- `make_figures.py` — reproduces the five figures in this report from `run22.log`.

B Reproducibility

The benchmark and report assume the four-directory layout described in Appendix A. From a working copy of that layout:

```
# Run the benchmark (~6 minutes wall clock):
cd ADASS2019 && python3 Run.py | tee ../logs/runN.log

# Update the frozen reference and regenerate figures:
cp ../logs/runN.log ../reports/run22.log
cd ../reports && python3 make_figures.py

# Recompile this document:
pdflatex Report-2026.tex && pdflatex Report-2026.tex
```

Output sorted by relative speed appears at the bottom of the run, after the per-test `Rept: ... Elap: ... 1K Iter:` lines and the `Fastest combination is ...` marker.

C Raw relative-time outputs

The complete sorted list of all 131 test variants from `run22.log` follows. The **Tier** column carries the same colour-coding as Figure 1 (green = fastest, purple = pathological); **Rel. time** is the geometric ratio to the hand-tuned arm64 NEON assembler (1.00); **ms / 1k iter** is the absolute time per 1000 iterations of the kernel, in milliseconds. This appendix is auto-generated from the canonical log by `make_raw_table.py`; running that script after replacing `run22.log` with a newer benchmark output will regenerate the table verbatim.

Tier	Rel. time	Family	Variant	ms / 1k iter
■	1.00	Assembler	clang (arm64)	0.0011
■	1.17	Julia	julia (int32)	0.0013
■	1.19	C : raw	g++ -O3 native	0.0013
■	1.20	C : raw	clang -O3 native	0.0013
■	1.21	C : Num. rec.	g++ -O3 native	0.0014
■	1.22	C : Num. rec.	clang -O3 native	0.0014
■	1.28	C : raw	g++ -O2	0.0014
■	1.28	C : raw	g++ -O3	0.0014
■	1.29	C : Num. rec.	g++ -O2	0.0014
■	1.30	C : Num. rec.	g++ -O3	0.0015
■	1.30	C : vectors	clang -O3 native	0.0015
■	1.30	C : vectors	g++ -O3 native	0.0015
■	1.37	C : vectors	g++ -O3	0.0015
■	1.40	C : vectors	g++ -O2	0.0016
■	1.41	C : Num. rec.	clang -O3	0.0016
■	1.41	C : raw	clang -O2	0.0016
■	1.49	C : raw	clang -O3	0.0017
■	1.49	C : vectors	clang -O3	0.0017
■	1.51	C : vectors	clang -O2	0.0017
■	1.54	C : Num. rec.	clang -O2	0.0017
■	2.05	Fortran	gfortran -O3 native	0.0023
■	2.50	Fortran	gfortran -O3	0.0028
■	2.51	Fortran	gfortran -O2	0.0028
■	2.59	Julia	julia -O2	0.0029
■	2.60	Julia	julia -O3	0.0029
■	2.76	D	ldc2 -O3 -release	0.0031
■	2.84	Zig	zig -O ReleaseFast	0.0032
■	3.15	Python : numpy raw	Python3 numba	0.0035
■	3.23	C++ : Boost no assert	g++ -O3 native	0.0036
■	3.24	C++ : Boost no assert	clang -O3 native	0.0036
■	3.25	C++ : Boost no assert	g++ -O2	0.0036
■	3.25	C++ : Boost no assert	g++ -O3	0.0036
■	3.31	Nim	nim -d:danger	0.0037
■	3.40	Julia	julia	0.0038
■	3.67	Crystal	crystal –release	0.0041
■	3.70	Java	Java	0.0042
■	4.30	Kotlin	kotlinc + JVM	0.0048
■	4.59	C : Num. rec.	clang -O1	0.0051
■	4.59	C : raw	g++ -O	0.0051
■	4.59	C : raw	g++ -O1	0.0051
■	4.59	C++ : Boost no assert	clang -O3	0.0051
■	4.60	C : raw	clang -O1	0.0052
■	4.62	C++ : Boost no assert	clang -O2	0.0052
■	4.64	Fortran	gfortran -O1	0.0052
■	4.65	C : vectors	clang -O1	0.0052
■	4.65	C : vectors	g++ -O	0.0052
■	4.66	C : vectors	g++ -O1	0.0052
■	4.69	C : Num. rec.	g++ -O1	0.0053
■	4.70	C : Num. rec.	g++ -O	0.0053
■	4.72	C : raw	clang -O	0.0053

continued on next page

(continued from previous page)

Tier	Rel. time	Family	Variant	ms / 1k iter
	4.73	C : Num. rec.	clang -O	0.0053
	4.78	C : vectors	clang -O	0.0054
	4.87	Fortran	gfortran -O	0.0055
	5.31	C#	.NET 8	0.006
	5.39	C#	.NET 8 AOT	0.006
	5.47	Perl/PDL : pp	Perl/PDL	0.0061
	5.50	Go	go -B (no BCE)	0.0062
	5.55	C++ : Boost no assert	g++ -O1	0.0062
	5.60	C++ : Boost no assert	g++ -O	0.0063
	6.76	Rust	Rustc iterators	0.0076
	6.78	C++ : Boost no assert	clang -O	0.0076
	6.91	Go	go	0.0078
	7.21	F#	.NET 8	0.0081
	7.59	Lua	LuaJIT	0.0085
	8.24	J	jconsole	0.0092
	8.48	Rust	Rustc unsafe	0.0095
	9.16	Haskell	GHC -O2	0.01
	9.44	OCaml	ocamlopt -O3	0.011
	10.23	C++ : Boost	clang -O3 native	0.011
	10.63	Lisp	SBCL typed	0.012
	10.84	C++ : Boost	g++ -O3 native	0.012
	11.39	C++ : Boost	g++ -O3	0.013
	12.04	Rust	Rustc -O3	0.013
	12.08	Rust	Rustc -O3 native	0.014
	12.29	Javascript	Bun	0.014
	12.49	C++ : Boost	clang -O3	0.014
	13.27	Rust	Rustc -O	0.015
	14.60	C++ : Boost	g++ -O2	0.016
	15.03	C++ : Boost	g++ -O	0.017
	15.80	C++ : Boost	clang -O2	0.018
	15.98	C++ : Boost	clang -O	0.018
	16.10	Python : vectors	Python3	0.018
	17.80	Python : arrays	Python3	0.02
	17.91	Swift	Swiftc -Ounchecked	0.02
	18.14	Swift	Swiftc -O	0.02
	21.15	Ruby : Numo	ruby 3.3 +YJIT	0.024
	21.24	Fortran	gfortran	0.024
	21.54	Ruby : Numo	ruby 3.3	0.024
	22.50	C : raw	g++	0.025
	22.56	C : Num. rec.	g++	0.025
	23.80	C : raw	clang	0.027
	23.89	C : Num. rec.	clang	0.027
	24.43	Javascript	Deno	0.027
	25.08	Perl/PDL : arrays	Perl/PDL	0.028
	25.88	Javascript	Node.js	0.029
	26.54	Javascript	Node.js –no-opt	0.03
	31.35	Python : vectors	Python2	0.035
	33.87	C++ : Boost	clang -O1	0.038
	36.17	Python : lists	PyPy3	0.041
	38.02	C++ : Boost no assert	clang -O1	0.043
	39.05	Python : arrays	Python2	0.044
	46.72	Julia	julia -O1	0.052
	47.17	C++ : Boost	g++ -O1	0.053
	54.60	C : vectors	g++	0.061
	55.72	C : vectors	clang	0.062
	57.45	Octave	octave vectorised	0.064
	66.85	Perl/PDL : vectors	Perl/PDL	0.075
	85.57	Julia	julia -O0	0.096
	150	R : outer	Rscript	0.17
	167	Lisp	SBCL untyped	0.19

continued on next page

(continued from previous page)

Tier	Rel. time	Family	Variant	ms / 1k iter
■	284	Ruby : raw	ruby 3.3 +YJIT	0.32
■	418	Rust	Rustc	0.47
■	644	Python : lists	Python3	0.72
■	716	Perl	Perl	0.8
■	755	Ruby : raw	ruby 3.3	0.85
■	782	Ruby : raw	ruby 2.6	0.88
■	794	R : raw	Rscript	0.89
■	878	C++ : Boost no assert	g++	0.98
■	905	C++ : Boost no assert	clang	1
■	918	Python : lists	Python2	1
■	921	C++ : Boost	g++	1
■	1012	C++ : Boost	clang	1.1
■	2883	Python : numpy raw	Python3	3.2
■	3599	Swift	Swiftc	4
■	3680	Swift	Swiftc -Onone	4.1
■	7161	Python : numpy raw	Python2	8
■	9207	Perl/PDL : raw	Perl/PDL	10
■	13,456	Tcl : raw	tclsh	15
■	22,015	Python : numpy raw	PyPy3	25
■	46,705	Octave	octave loop	52
■	481,009	R : ref class	Rscript	5.4e+02